

"Equal Opportunity" and Diversity Compliance: AI/ML Job Postings in a Post-DEI America

Evelyn Yee

Department of Computer Science
Stanford University
yeevelyn@stanford.edu

Abstract

With the widespread impact that artificial intelligence and machine learning (AI/ML) technologies have had on all aspects of society, it more important than ever to ensure that diverse perspectives are involved in the development of this transformative technology. To understand how these tech companies approach diversity in their hiring practices, I analyzed text features of AI/ML job postings from 2022 to May 2025. Using dictionary-based metrics and binary annotations, I categorize job postings by their DEI sentiment, finding that DEI language in AI/ML postings seems to be heavily motivated by regulatory compliance, shifting accordingly in response to the anti-DEI executive orders that Donald Trump enacted, starting at his inauguration in January 2025.¹

1 Introduction

Generative AI technology is having an increasing impact on the social, political, and economic conditions in the world (Maslej et al., 2025; Hagerty and Rubinov, 2019; Bommasani, 2025). Leading this wave of impacts changes are a few major U.S.-based technology companies, such as OpenAI and Anthropic. As AI models increasingly shape how people access information (Mo et al., 2025), make decisions (Gorka et al., 2025), and perform work (Handa et al., 2025), it is especially important to consider biases and harms within these models. In particular, model design choices are shaped by stakeholder involvement in the AI development process, principally through employee perspectives. From a professional front, integrating diverse perspectives increases safety compliance (Lacerenza et al., 2025), and in technological contexts, it supports innovation (Østergaard et al., 2011). For these reasons, the current work sheds insight on em-

ployee diversity measures in AI companies through the lens of job postings.

Job postings serve as the primary point of entry for many prospective employees, informing their decision to apply to a job. So, companies have an incentive to represent their working culture as welcoming, perhaps pro-DEI. However, with the recent political and social shift away from progressive movements like DEI and antiracism (Wetts and Willer, 2025), I investigated whether public-facing language in job postings has shifted as well, perhaps creating a more hostile appearance for diverse applicants.

In the United States, there is a standing precedent of regulations around non-discrimination and diversity in hiring practices. Nearly all companies in the U.S. are bound by the Civil Rights Act of 1964, formally barring companies from discriminating against job applicants on the basis of "race, color, religion, or national origin" (U.S. Congress, 1964). The Equal Employment Opportunity Commission monitors and adjudicates companies' adherence to this law. Additionally, many state-level and local jurisdictions have enacted further laws guiding anti-discrimination in hiring, such as the San Francisco Fair Chance Ordinance, which restricts employers from inquiring about applicants' arrest or conviction records (California State Assembly, 2017).

In January 2025, however, the Trump administration issued a series of executive orders targeting diversity, equity, and inclusion (DEI) initiatives across federal agencies and federal contractors, framing such programs as "radical and wasteful," causing "illegal discrimination" (Executive Office of the President, 2025a,b). This political shift was further reinforced by the administration's AI Action Plan, which explicitly positioned the rejection of "woke" principles as a cornerstone of U.S. AI policy (The White House, 2025). While these executive orders have direct legal force only over government entities and certain contractors, they

¹Data and code for this project are available at https://github.com/evelynyee/jobs_dei

signal a broader political environment that may influence corporate behavior, particularly among technology companies seeking to navigate an increasingly polarized landscape.

In this work, I analyze the prevalence of overt DEI sentiment in AI/ML domain job postings from the last 3 years, as quantified through dictionary-based methods and holistic, whole-text annotations. After separating out job postings that contain DEI sentiment from those that do not, I inspect the descriptive qualities of these two groupings, regarding the companies that make the posts and the specific language used. Finally, I interpret these findings under a regulatory framework, finding evidence that tech companies use DEI language in their job postings as a compliance method to meet legal obligations.

2 Related Works

2.1 Biased Language in Job Postings

Diversity measures in the text of job postings have been primarily studied in regard to gender bias. For example, many works use NLP methods on job postings to evaluate the presence of gender discriminatory sentiment (Chaturvedi et al., 2021; Mihaljević et al., 2022; Ningrum et al., 2020), the intersections of gender bias with other kinds of bias, such as racial bias and age bias (Hu et al., 2024; Frissen et al., 2023; Burn et al., 2022). Further works suggest tools to mitigate gender bias in job postings (Hu et al., 2022; He and Kang, 2025; Horvath and Sczesny, 2016), and some works have investigated the downstream effects of gender bias on job applicants and the labor force (Squicciarini et al., 2024; Chaturvedi et al., 2021). There are many works investigating gender bias in job postings, of which I have cited a selection.

Additionally, a range of works study the expression and impact of age-related bias in job postings, especially within technical fields (Burn et al., 2019; Baltes et al., 2020).

2.2 Biased Language in Other Contexts

Additional works have explored socially exclusive and "biased" language in other contexts. For example, Chakravarthi (2024) and Baltes et al. (2020) investigated bias in online social media posts (relating to homophobia/transphobia and age bias, respectively), using manual annotations to identify bias. Additionally, many works have used NLP methods to evaluate gender bias in text (Zhu et al.,

2024; Orgeira-Crespo et al., 2020) and mitigate its impact in a variety of contexts (Todd et al., 2024). Most similarly to the current work, Chakravarthi et al. (2023) classified marketing communications according to four types of bias (age, disability, gender, and "other"), using dictionary methods and annotation. My methodology was heavily inspired by their process for bias classification, using word embeddings to build a dictionary and comparing dictionary heuristics to holistic document annotations.

3 Methods

3.1 Datasets




Figure 1: Distribution of job postings within the two datasets.

I began by using two free, public job posting datasets from Kaggle. The first dataset (Gajamuthu, 2024) collects LinkedIn job postings from 2022-2024, claiming that it "provides a comprehensive collection of job postings in the field of Artificial Intelligence (AI) and Machine Learning (ML) across the United States." The second dataset (Kumeyko, 2025) contains job postings from "Machine Learning-related roles across the United States," posted on a variety of job boards between 2022 and April 2025. Each dataset has approximately 1000 job postings.

The datasets show a peak in job postings around May 2024 and April 2025, which could indicate a hiring cycle in the spring, but it is most likely that these datasets were just scraped in the spring of their respective years (i.e. because more recent job postings are more likely to still be accessible online).

Inspired by Chakravarthi et al. (2023)'s analytical process, I quantified DEI sentiment in these job postings in two phases: using dictionary-based word heuristics and then holistic annotations.

3.2 Dictionary Methods

To identify DEI sentiment, the first approach I took was to count the occurrences of DEI-signaling words in job postings. Specifically, I looked for words that would explicitly indicate that principles of diversity, equity, and inclusion were company priorities. While further work should be done to identify more subtle, covert indicators of diversity-friendly hiring practices, I chose to focus on public-facing language that was unambiguously related to diversity, equity, and inclusion, as well as related initiatives like nondiscrimination practices.

On a surface level, companies can leverage these DEI buzzwords for shallow virtue signaling purposes, invoking a public perception of virtue without actually making substantive diversity efforts. (Berthon et al., 2023) However, I find that even if companies are using DEI language for virtue signaling purposes, that action indicates that the company expects some cultural or financial incentive for advertising diversity, which is related to my objective of measuring DEI sentiment in the AI hiring pipeline.

3.2.1 Text Normalization

For the dictionary-based methods, I normalized all of the job posting texts by converting characters to lowercase, removing punctuation, and combining whitespaces. I then tokenized the job descriptions by splitting between words. There were some normalization issues in the datasets (e.g. "benefitscompetitive"), likely due to artifacts of the data scraping process, but I believe that these occurrences were rare within the whole collection of job posting texts, and I also manually filtered these words out from my DEI vocabulary, as described below.

3.2.2 Building a DEI Vocabulary

I started by constructing a narrow vocabulary of 7 DEI words: "dei", "diversity", "diverse", "equity", "equitable", "inclusion", and "inclusive."

To identify additionally DEI-related words, I trained a Word2Vec model (Mikolov et al., 2013) on this corpus of normalized job descriptions. I then used the learned embeddings to identify the 50 most similar words for each of the narrow DEI vocabulary words, filtering to only words with similarity over 0.6.² Due to the non-DEI, financial

²These similarity thresholds were tuned empirically to optimize size and dictionary quality over manual inspection of the resulting dictionaries.

meaning of "equity" in some job postings (i.e. ownership in a company), I also removed DEI terms that had similarity over 0.5 to any other term in a financial dictionary I came up with.³ Finally, I performed a manual filtering step, removing any terms that seemed to be stray artifacts of the text, such as "dibe" and "pluspaypal."

After filtering and accounting for repeats, this gave me an expanded vocabulary of 69 words. Some of these words seemed clearly related to diversity, equity, and inclusion, such as "ethical" and "culture," but some others were more ambiguous, such as "strive" and "gritty." (See Appendix A for the full set of words in the expanded DEI vocabulary.)

3.3 Holistic Annotations

To assess more subtle expressions of DEI sentiment, I annotated each complete text with an LLM, OpenAI's GPT-5 Nano. For each job posting, the LLM was prompted to respond with a binary score, corresponding to whether the posting contained DEI sentiment or not. (See Appendix B for the full prompt text.)

I validated the LLM annotations over a random sample of 50 manually annotated job postings, finding substantial agreement between the LLM and my judgments (Cohen's $\kappa = 0.8$).

4 Results

4.1 Dictionary Methods




Figure 2: Frequency of DEI words (for both the narrow and expanded DEI vocabularies) in each job posting, as a percent of the posting length, averaged per month. Lines of best-fit are additionally plotted, where "*" indicates $p < 0.05$, "***" indicates $p < 0.01$, "****" indicates $p < 0.001$.

³Interestingly, this also filtered out the word "dei," which had high similarity to "profit."

To estimate the prevalence of DEI content in job postings, I counted the frequency of my identified vocabulary words within the job posting texts over time. I observed a high level of fluctuation, with a significant but weak upward trend over time. The low r^2 values on these best-fit lines indicate that there may be other more influential determiners of DEI language than time.

Plotting this against the day that Trump took office and started signing in anti-DEI executive orders, it seems that there may have been a dip in DEI language leading up to the presidential turnover and immediately after the executive orders. However, the levels of DEI language seem to have rebounded by March, matching the general trend over time.

4.2 Holistic Annotations




Figure 3: Proportion of job postings each month that expressed "DEI sentiment," as annotated by GPT-5 Nano. The line of best-fit is additionally plotted, where "*" indicates $p < 0.05$, "***" indicates $p < 0.01$, "****" indicates $p < 0.001$.

As seen in Figure 3, looking at DEI sentiment as a binary feature of job postings, I observed strikingly similar patterns between the holistic annotations as the previous word heuristics. The proportion of job postings expressing DEI sentiment dropped in January and February of 2025, before rebounding back to the best-fit line by March 2025. As with the dictionary word frequencies, the DEI annotations also showed a low r^2 value with respect to time, further confirming the weak relationship between time and DEI sentiment.

As seen in Figure 4, the annotated binary DEI label for a document seems to correlate with the proportion of words in that document matching our DEI vocabulary, both for the narrow, 7-word vocabulary ($r_{pb} = 0.425$, $p < 0.001$) and the expanded, 69-word vocabulary ($r_{pb} = 0.533$, $p < 0.001$). This establishes convergent validity between these




Figure 4: Distributions of narrow and expanded DEI word proportions, between postings that were annotated as expressing DEI sentiment vs. not expressing DEI sentiment.

two measures, addressing the common construct of DEI sentiment.

4.3 Descriptive Analysis

After using these two methods to classify job postings as DEI and non-DEI, I performed a variety of descriptive analyses to understand the defining characteristics of each posting type. See Appendix C for representative examples of DEI and non-DEI job postings.

4.3.1 Companies

Looking at specific companies, it is difficult to identify the trends over time because approximately 34% of the companies represented in this dataset had only one job posting to analyze. Focusing on the impact of the executive orders, only 9% of the companies in the dataset posted at least once before January 20, 2025 and at least once after. Because of this lack of resolution on the individual company level, I primarily viewed companies in the aggregate, looking for shared characteristics.

In particular, I noticed that the companies that were most prolific in their DEI posting were mostly large tech companies, such as TikTok (89 DEI postings) and Meta (37 DEI postings). These companies also rarely published non-DEI postings, as only 1 of the top 10 DEI posting companies also published any non-DEI postings.

In contrast, the most prolific posters of non-DEI job postings were specialized talent recruiting agencies, such as Unreal Staffing (35 non-DEI postings) and comparatively smaller startups, such as Waymo (10 non-DEI postings). The non-DEI posting companies were generally much less prolific in the dataset, and 3 of the top 10 non-DEI posting companies also published at least one DEI posting.

4.3.2 Language Use

To understand the specific language used in DEI vs. non-DEI job posts, I treated each category of postings as a corpus. Within each corpus, I counted the frequency of each word and compared the most common words from each. Looking at the top 400 words from each corpus, 285 of them were frequent in both corpora, such as "machine," "learning," "data," and "model."

However, of the 125 remaining uniquely common words for each corpus, I found a large difference between the language of DEI and non-DEI job postings. As seen in Figure 5, DEI job postings heavily referenced features of the employees and applicants (e.g., "disability" and "gender") as well as the employers' obligations (e.g., "accommodation" and "law"). In contrast, as seen in Figure 6, the non-DEI postings used more impersonal language, emphasizing technical tools (e.g., "azure" (Microsoft Azure) and "architecture"), work outputs (e.g., "client" and "deploying"), and judgments of quality (e.g. "excellent" and "enhance").

Specifically, I noticed that over 70% of the DEI job postings specifically referenced their Equal Employment Opportunity (EEO) obligation (U.S. Congress, 1964), whereas only one non-DEI job posting in the dataset mentioned EEO.



Figure 5: The 125 most frequent words in DEI postings that were not frequent in non-DEI postings. Size indicates word frequency within the DEI corpus.

5 Discussion

Viewing these results, I hypothesize that these AI companies deploy DEI language in their job postings to meet federal, state, and local regulations. Under this framework, I explain three major findings from this analysis: the DEI response to the January 2025 executive orders, the company discrepancies between DEI and non-DEI postings, and the descriptive differences between DEI and non-




Figure 6: The 125 most frequent words in non-DEI postings that were not frequent in DEI postings. Size indicates word frequency within the non-DEI corpus.

DEI language use.

5.1 Regulatory Anticipation

In January 2025, the Trump administration burst through with sweeping promises to overhaul regulation, starting with the many executive orders they enacted over the first 2 days of the term, including those declaring DEI programs as a type of discrimination. Concurrently, the administration expounded a flood of related rhetoric. From the perspective of a tech company at this time, these executive orders () may forecast a shift in federal DEI regulation, forcing companies to abandon their public DEI-favorable hiring practices. In my numerical results, this effect would explain the sharp decrease in DEI sentiment within AI/ML job postings during the time period leading up to Trump's inauguration and immediately after he signed in the anti-DEI executive orders. The initial decline in DEI language suggests that companies responded preemptively in anticipation of future regulatory changes.

However, although these executive orders significantly impacted programs within the federal government and produced some additional paperwork for federal government contractors, they did not directly impact the bulk of work performed by private sector tech companies, which hire for AI and ML jobs. In fact, the federal, state, and local regulations requiring nondiscrimination in private companies' hiring (and encouraging or requiring disclosure of these commitments in job postings) had not been modified by the executive orders. After the initial commotion settled, tech companies would have realized that their compliance obligations regarding diversity and nondiscrimination, particularly in relation to the Equal Employment Opportunity Commission, had not significantly changed. This

hypothesis aligns with the DEI language rebound I observed in job postings during March and April of 2025, as tech companies returned to the status quo.

This event-centered regulatory interpretation is reinforced by the low r^2 values for the linear fit lines between time and both the dictionary-based and holistic annotation methods; if DEI language were driven by genuine cultural evolution in organizational values, we would expect smoother, more consistent trends over time. Instead, the volatility and weak time trends suggest that, within the time scale of this data, DEI language in job postings responded more strongly to discrete events, such as legal guidance updates and risk assessments, rather than larger-scale shifting of cultural norms.

5.2 Corporate Compliance Structures

This regulatory approach also aligns with my descriptive findings around job posting companies. Specifically, large, established technology companies face heightened regulatory scrutiny due to their size, public visibility, and frequent interaction with government entities. This could explain why the largest companies (the most prolific posters in the dataset) heavily favored DEI sentiment in their job postings. Additionally, such large companies have the organizational infrastructure to have standardized compliance protocols, leading to uniform templates for job postings, as demonstrated by the fact that the top DEI posting companies produced almost no deviant, non-DEI job postings.

In contrast, recruiting agencies may operate under different regulatory frameworks or may not feel compelled to include DEI language when posting on behalf of clients, as the legal responsibility for hiring compliance ultimately lies with the employer rather than the intermediary. Similarly, smaller startups, while still legally required to comply with EEO regulations, may face lower practical enforcement risk due to reduced visibility, fewer regulatory touchpoints, and more limited resources for legal oversight. This would explain why the more prolific non-DEI posting organizations were all companies of these types. Furthermore, the variability in these companies' posting behaviors, where multiple of the top non-DEI posters also published some DEI job postings, may reflect more discretionary or situational compliance systems within these organizations.

5.3 Language of Compliance

Contrasting the use of language in DEI and non-DEI job postings, we can see clear evidence of regulatory impact on the wording of these postings. Most obviously, the words "equal," "employment," and "employer" were among the most frequent words in DEI postings but not non-DEI posts. This is a clear reference to the "equal employment opportunity" (EEO) obligations that almost every American company has as an "equal opportunity employer."⁴ The fact that EEO mentions appear almost exclusively in DEI-classified postings suggests that the presence of DEI language is closely tied to legal disclosure practices mandated or encouraged by civil rights legislation (U.S. Congress, 1964).

Additionally, the concentration in DEI job postings of specific terms referencing legally protected classes, such as "disability," "gender," "race," and "national" "origin," indicates that companies are not simply expressing general values about diversity but are specifically addressing their legal responsibilities (i.e., their "commitment" to the "law"), especially in relation to workplace "accommodation" and how employees would be "paid." This pattern stands in sharp contrast to the language of non-DEI postings, which impersonally focus on technical capabilities and business outcomes without reference to employment law or worker characteristics. This divergence between the heavy emphasis on legally relevant language in DEI job postings and the complete absence of personal characteristics in non-DEI job postings suggests that tech companies treat DEI content as a distinct legal addendum to their job postings rather than a key aspect of connecting with applicants in the hiring process.

6 Limitations

The primary limitation of this work is in its scale. Due to data availability constraints, I was only able to find these two high-quality, publicly available datasets for American AI/ML job postings. As a result, my quantitative results were quite noisy, and my conclusive power is very limited, especially for making clear before-and-after comparisons, considering that certain time periods were over- and under-represented in this dataset (i.e., not spring

⁴The word "opportunity" was the 45th most common word in the DEI postings, but it was also 203rd most common in the non-DEI postings.

2024 or spring 2025). Additionally, these qualities of my data also limit my ability to draw a clear causal relationship between the executive orders and the job postings (and to rule out other possible causes for the DEI fluctuation, such as economic factors); I can only make descriptive observations.

Additionally, the scope of this analysis and conclusions were very specific to this small career domain of AI and machine learning, within the US technology ecosystem. To get a more robust view of the impact of the January 2025 executive orders on job posting trends, it would be important to evaluate job postings from other industries, as well as public sector and non-profit organizations.

7 Conclusion

This work analyzes the prevalence of anti-discrimination and DEI-related language in American job postings relating to artificial intelligence and machine learning. I found that regulatory requirements (and the anticipation of such) seem to correspond to companies' use of DEI language in these roles, indicating the power of regulation in this space. However, further work is needed to identify the downstream impacts of DEI language and similar observable diversity practices on actual hiring outcomes in tech companies and the specific impact of employee diversity on resultant AI systems.

Seeing the empirical impact of regulation on this one aspect of diversity in tech hiring, these results encourage future work into other ways that regulatory mechanisms can be shaped to more deeply integrate these effects into AI/ML hiring processes and the broader technology development lifecycle. However, the current political discourse in the United States is swinging sharply in an anti-diversity direction, especially on a federal level. With this in mind, I find particular room for works investigating the regulatory motives of state and local governments as well as the role of internal industry norms and individual applicants/employees to self-encourage diverse hiring practices, as predicted by Fan (2019).

References

Sebastian Baltes, George Park, and Alexander Serebrenik. 2020. [Is 40 the new 60? how popular media portrays the employability of older software developers](#). *Preprint*, arXiv:2004.05847.

Pierre R. Berthon, Sarah T. Lord Ferguson, Leyland F.

Pitt, and Emma Wang. 2023. [The virtuous brand: The perils and promises of brand virtue signaling](#). *Business Horizons*, 66(1):27–36.

Rishi Bommasani. 2025. [The societal impact of foundation models: Advancing evidence-based ai policy](#). *Preprint*, arXiv:2506.23123.

Ian Burn, Patrick Button, Luis Munguia Corella, and David Neumark. 2022. [Does ageist language in job ads predict age discrimination in hiring?](#) *Journal of Labor Economics*, 40(3):613–667.

Ian Burn, Patrick Button, Luis Felipe Munguia Corella, and David Neumark. 2019. [Older workers need not apply? ageist language in job ads and age discrimination in hiring](#). Working Paper 26552, National Bureau of Economic Research.

California State Assembly. 2017. [An act to add section 12952 to the government code, and to repeal section 432.9 of the labor code, relating to employment discrimination](#). A.B. 1008, 2017-2018 Leg., Reg. Sess.

Bharathi Raja Chakravarthi. 2024. [Detection of homophobia and transphobia in youtube comments](#). *International Journal of Data Science and Analytics*, 18(1):49–68.

Bharathi Raja Chakravarthi, Prasanna Kumar Kumaresan, Rahul Ponnusamy, John P. McCrae, Michaela Comerford, Jay Megaro, Deniz Keles, and Last Feremenga. 2023. [Exploring techniques to detect and mitigate non-inclusive language bias in marketing communications using a dictionary-based approach](#). In *Proceedings of the 14th International Conference on Recent Advances in Natural Language Processing*, pages 918–925, Varna, Bulgaria. INCOMA Ltd., Shoumen, Bulgaria.

Sugat Chaturvedi, Kanika Mahajan, and Zahra Siddique. 2021. [Words matter: Gender, jobs and applicant behavior](#). IZA Discussion Paper 14497, Institute of Labor Economics (IZA).

Executive Office of the President. 2025a. [Executive order 14151: Ending radical and wasteful government dei programs and preferencing](#). Federal Register.

Executive Office of the President. 2025b. [Executive order 14173: Ending illegal discrimination and restoring merit-based opportunity](#). Federal Register.

Jennifer S. Fan. 2019. [Employees as regulators: The new private ordering in high technology companies](#). *Utah Law Review*, 2019(5):973–1026.

Richard Frissen, Kolawole John Adebayo, and Rohan Nanda. 2023. [A machine learning approach to recognize bias and discrimination in job advertisements](#). *AI & SOCIETY*, 38(2):1025–1038.

Kanchana Gajamuthu. 2024. [AI and ML job listings USA](#).

- Ernest Gorka, Dariusz Baran, Gabriela Wojak, Michal Ćwiakala, Sebastian Zupok, Dariusz Starkowski, Dariusz Resko, and Oliwia Okrasa. 2025. [The impact of artificial intelligence on enterprise decision-making processes](#). *EUROPEAN RESEARCH STUDIES JOURNAL*, XXVIII(Issue 4):756–777.
- Alexa Hagerty and Igor Rubinov. 2019. [Global ai ethics: A review of the social impacts and ethical implications of artificial intelligence](#). *Preprint*, arXiv:1907.07892.
- Kunal Handa, Alex Tamkin, Miles McCain, Saffron Huang, Esin Durmus, Sarah Heck, Jared Mueller, Jerry Hong, Stuart Ritchie, Tim Belonax, Kevin K. Troy, Dario Amodei, Jared Kaplan, Jack Clark, and Deep Ganguli. 2025. [Which economic tasks are performed with ai? evidence from millions of claude conversations](#). *Preprint*, arXiv:2503.04761.
- Joyce C. He and Sonia K. Kang. 2025. [Debiasing job ads by replacing masculine language increases gender diversity of applicant pools](#). *Proceedings of the National Academy of Sciences*, 122(7):e2409854122.
- Lisa Kristina Horvath and Sabine Sczesny. 2016. [Reducing women’s lack of fit with leadership positions? effects of the wording of job advertisements](#). *European Journal of Work and Organizational Psychology*, 25(2):316–328.
- Shenggang Hu, Jabir Alshehabi Al-Ani, Karen D. Hughes, Nicole Denier, Alla Konnikov, Lei Ding, Jinhua Xie, Yang Hu, Monideepa Tarafdar, Bei Jiang, Linglong Kong, and Hongsheng Dai. 2022. [Balancing gender bias in job advertisements with text-level bias mitigation](#). *Frontiers in Big Data*, 5:805713.
- Yang Hu, Nicole Denier, Lei Ding, Monideepa Tarafdar, Alla Konnikov, Karen D Hughes, Shenggang Hu, Bran Knowles, Enze Shi, Jabir Alshehabi Al-Ani, Irina Rets, Linglong Kong, Dengdeng Yu, Hongsheng Dai, and Bei Jiang. 2024. [Language in job advertisements and the reproduction of labor force gender and racial segregation](#). *PNAS Nexus*, 3(12):pgae526.
- Ivan Kumeiko. 2025. [Machine learning job postings in the US](#).
- Christina N. Lacerenza, Stefanie K. Johnson, Natalie V. Schwatka, Marissa A. Beldon, and Jack T. Dennerlein. 2025. [Team diversity as a safety asset: A field investigation of language diversity and occupational safety](#). *Human Resource Management*, 64(3):695–711.
- Nestor Maslej, Loredana Fattorini, Raymond Perrault, Yolanda Gil, Vanessa Parli, Njenga Kariuki, Emily Capstick, Anka Reuel, Erik Brynjolfsson, John Etchemendy, Katrina Ligett, Terah Lyons, James Manyika, Juan Carlos Niebles, Yoav Shoham, Russell Wald, Toby Walsh, Armin Hamrah, Lapo Santarasci, and 4 others. 2025. [The ai index 2025 annual report](#). Technical report, AI Index Steering Committee, Institute for Human-Centered AI, Stanford University, Stanford, CA.
- Helena Mihaljević, Ivana Müller, Katja Dill, and Aysel Yollu-Tok. 2022. [Towards gender-inclusive job postings: A data-driven comparison of augmented writing technologies](#). *PLOS ONE*, 17(9):1–35.
- Tomas Mikolov, Kai Chen, Greg Corrado, and Jeffrey Dean. 2013. [Efficient estimation of word representations in vector space](#). *Preprint*, arXiv:1301.3781.
- Fengran Mo, Kelong Mao, Ziliang Zhao, Hongjin Qian, Haonan Chen, Yiruo Cheng, Xiaoxi Li, Yutao Zhu, Zhicheng Dou, and Jian-Yun Nie. 2025. [A survey of conversational search](#). *Preprint*, arXiv:2410.15576.
- Panggih Kusuma Ningrum, Tatdow Pansombut, and Attachai Ueranantasun. 2020. [Text mining of online job advertisements to identify direct discrimination during job hunting process: A case study in indonesia](#). *PLOS ONE*, 15(6):1–29.
- Pedro Orgeira-Crespo, Carla Míguez-Álvarez, Miguel Cuevas-Alonso, and María Isabel Doval-Ruiz. 2020. [Decision algorithm for the automatic determination of the use of non-inclusive terms in academic texts](#). *Publications*, 8(3).
- Mariagrazia Squicciarini, Layla O’Kane, Julia Nania, Echo Liu, and Scott Bingham. 2024. [The weight of words: Gendered language and women’s participation and positioning in the labor market](#). MOST Policy Papers 3, UNESCO, Paris, France. Foreword by Gabriela Ramos.
- The White House. 2025. [Winning the race: America’s AI action plan](#).
- Liam Todd, John Grundy, and Christoph Treude. 2024. [Githubinclusifier: Finding and fixing non-inclusive language in github repositories](#). In *Proceedings of the 2024 IEEE/ACM 46th International Conference on Software Engineering: Companion Proceedings, ICSE-Companion ’24*, page 89–93, New York, NY, USA. Association for Computing Machinery.
- U.S. Congress. 1964. [An act to enforce the constitutional right to vote, to confer jurisdiction upon the district courts of the united states of america to provide injunctive relief against discrimination in public accommodations, to authorize the attorney general to institute suits to protect constitutional rights in public facilities and public education, to extend the commission on civil rights, to prevent discrimination in federally assisted programs, to establish a commission on equal employment opportunity, and for other purposes](#). H.R. 7152, 88th Cong. Enacted July 2, 1964.
- Rachel Wetts Wetts and Robb Willer. 2025. [Antiracism and its discontents: Opposition to antiracism is a widespread and politically influential racial ideology among white americans](#).

Haotian Zhu, Kexin Gao, Fei Xia, and Mari Ostendorf. 2024. [Disagreeable, slovenly, honest and un-named women? investigating gender bias in English educational resources by extending existing gender bias taxonomies](#). In *Proceedings of the 5th Workshop on Gender Bias in Natural Language Processing (GeBNLP)*, pages 219–236, Bangkok, Thailand. Association for Computational Linguistics.

Christian R. Østergaard, Bram Timmermans, and Kari Kristinsson. 2011. [Does a different view create something new? the effect of employee diversity on innovation](#). *Research Policy*, 40(3):500–509.

A DEI Vocabulary Details

A.1 Negative Financial Vocabulary

To filter out the financial meanings associated with "equity," I removed words with similarity to any of the following:

"economic", "expense", "expensive", "finance", "financial", "fiscal", "pay", "profit", "salary"

A.2 Expanded DEI Vocabulary

After the financial filtering and a final process of manual filtering, the resulting expanded DEI vocabulary words are provided here in alphabetical order:

"act", "adaptable", "aim", "amazing", "atmosphere", "backgrounds", "belonging", "belongs", "celebrate", "celebrates", "celebrating", "collaborative", "commitment", "committed", "communicative", "communities", "competitive", "creating", "cultivating", "culture", "dedication", "differences", "diverse", "diversity", "embrace", "embraces", "embracing", "empathetic", "employer", "empowering", "encourages", "equal", "equality", "equitable", "ethical", "foster", "fostering", "free", "fun", "gives", "grant", "gritty", "happiness", "honesty", "inclusion", "inclusive", "inclusivity", "inspired", "intentionally", "personify", "promote", "providing", "remarkable", "respect", "respectful", "retaining", "rewarding", "society", "spirit", "stayed", "strive", "stronger", "supportive", "values", "viewpoints", "voices", "welcoming", "workforce", "workplace"

B LLM Annotation Details

B.1 Annotation Prompt

System Prompt:

You are a text processing assistant for a recruiting and human resources company, helping them process job postings. They are classifying job postings according to a few characteristics. Each user

message contains the description of a single characteristic and a job posting text. You will respond with the classification of the job posting according to the description.

User Prompt:

CHARACTERISTIC: Diversity, Equity, and Inclusion (DEI) Sentiment

DESCRIPTION: The job posting contains information encouraging diversity in applicants. For example, mention of a welcoming workplace, emphasis on diverse perspectives, and statement of inclusive policies, such as affirmative action and Equal Employment Opportunity/being an Equal Opportunity Employer.

CATEGORIZATION: 1 (includes DEI sentiment), 0 (does not include DEI sentiment)

JOB POSTING:

""

{text}

""

Think step-by-step and then respond with "CATEGORIZATION:" followed by your categorization (only 0 or 1).

B.2 Validation Confusion Matrix

	LLM: No DEI	LLM: DEI	Total
Human: No DEI	23	3	26
Human: DEI	2	22	24
Total	25	25	50

C Representative Job Postings

C.1 DEI Job Posting: TikTok

Responsibilities

TikTok is the leading destination for short-form mobile video. Our mission is to inspire creativity and bring joy. TikTok has global offices including Los Angeles, Mountain View, New York, Austin, London, Paris, Berlin, Dubai, Singapore, Jakarta, São Paulo, Seoul and Tokyo.

Why Join Us

Creation is the core of TikTok's purpose. Our platform is built to help imaginations thrive. This is doubly true of the teams that make TikTok possible. Together, we inspire creativity and bring joy - a mission we all believe in and aim towards achieving every day. To us, every challenge, no matter how difficult, is an opportunity; to learn, to innovate, and to grow as one team. Status quo? Never. Courage? Always. At TikTok, we create together

and grow together. That's how we drive impact - for ourselves, our company, and the communities we serve.

Team Introduction:

Join the E-commerce Global Supply Chain and Logistics team at TikTok! We're enhancing the shopping experience and reducing logistics operational costs. We are seeking brilliant and motivated graduate software engineers, who are eager to apply their knowledge in machine learning (ML), operations research (OR), data mining, and statistical inference to real-world challenges.

We are looking for talented individuals to join our team in 2024. As a graduate, you will get unparalleled opportunities for you to kickstart your career, pursue bold ideas and explore limitless growth opportunities. Co-create a future driven by your inspiration with TikTok.

Successful candidates must be able to commit to one of the following start dates below:

1. January 15, 2024
2. February 5, 2024
3. March 4, 2024
4. May 20, 2024
5. June 10, 2024
6. July 15, 2024
7. August 12, 2024

We will prioritize candidates who are able to commit to these start dates. Please state your availability and graduation date clearly in your resume.

Applications will be reviewed on a rolling basis. We encourage you to apply early.

Candidates can apply for a maximum of TWO positions and will be considered for jobs in the order you applied for. The application limit is applicable to TikTok and its affiliates globally.

Join our vibrant and fast-paced team at TikTok as a Machine Learning Engineer, specializing in Computer Vision (CV), Artificial Intelligence Generated Content (AIGC), or Multimodal machine learning. This role offers an exciting opportunity for recent graduates (Master's or PhD) who can start work in 2024 to apply state-of-the-art techniques to challenging problems in the e-commerce supply chain and logistics domain. Be part of the team revolutionizing customer experience through advanced technologies.

Responsibilities:

- Explore and implement cutting-edge technologies like computer vision, AIGC, and multimodal machine learning in content commerce.
- Leverage AIGC technology to enhance various aspects of e-commerce, including merchandising, product image optimization, and intelligent design.
- Optimize and iterate CV/multimodal/AIGC-related models, focusing on image understanding, generation, and editing.

Qualifications

- Master's degree in Computer Science, Engineering, Operations Research, or related fields.
- Graduate with a background in Software Development, Computer Science, Computer Engineering, or a related technical discipline
- Understanding of TensorFlow/PyTorch or similar training frameworks. Familiarity with mixed precision training, distributed training, and TensorRT deployment. Preferred Qualifications:
- Demonstrated software engineering experience from previous internship, work experience, coding competitions, or publications
- Internship experience or research experience, especially in e-commerce, multimodal/AIGC, including image and video understanding, generation.
- Research in a specific area of computer vision/multimodality and publications at conferences such as KDD, NeurIPS, WWW, SIGIR, WSDM, CIKM, ICLR, ICML, IJCAI, AAAI.
- Excellent communication skills, strong learning ability, proactive, and enthusiastic attitude.

TikTok is committed to creating an inclusive space where employees are valued for their skills, experiences, and unique perspectives. Our platform connects people from across the globe and so does our workplace. At TikTok, our mission is to inspire creativity and bring joy. To achieve that goal, we are committed to celebrating our diverse voices and to creating an environment that reflects the many communities we reach. We are passionate about this and hope you are too.

TikTok is committed to providing reasonable accommodations in our recruitment processes for candidates with disabilities, pregnancy, sincerely held religious beliefs or other reasons protected by applicable laws. If you need assistance or a reasonable accommodation, please reach out to us at dataecommerce.accommodations@tiktok.com.

By submitting an application for this role, you accept and agree to our global applicant privacy policy, which may be accessed here: <https://careers.tiktok.com/legal/privacy>.

Job Information:

[For Pay Transparency] Compensation Description (annually)

The base salary range for this position in the selected city is \$106590 - \$139650 annually.

Compensation may vary outside of this range depending on a number of factors, including a candidate's qualifications, skills, competencies and experience, and location. Base pay is one part of the Total Package that is provided to compensate and recognize employees for their work, and this role may be eligible for additional discretionary bonuses/incentives, and restricted stock units.

Our company benefits are designed to convey company culture and values, to create an efficient and inspiring work environment, and to support our employees to give their best in both work and life. We offer the following benefits to eligible employees:

We cover 100% premium coverage for employee medical insurance, approximately 75% premium coverage for dependents and offer a Health Savings Account(HSA) with a company match. As well as Dental, Vision, Short/Long term Disability, Basic Life, Voluntary Life and AD&D insurance plans. In addition to Flexible Spending Account(FSA) Options like Health Care, Limited Purpose and Dependent Care.

Our time off and leave plans are: 10 paid holidays per year plus 17 days of Paid Personal Time Off (PPTO) (prorated upon hire and increased by tenure) and 10 paid sick days per year as well as 12 weeks of paid Parental leave and 8 weeks of paid Supplemental Disability.

We also provide generous benefits like mental and emotional health benefits through our EAP and Lyra. A 401K company match, gym and cellphone service reimbursements. The Company reserves the right to modify or change these benefits programs at any time, with or without notice.

C.2 Non-DEI Job Posting: Unreal Staffing

We're seeking exceptional engineers to drive our ambitious goals forward! We need individuals with extraordinary skills who can thrive in our fast-paced environment and deliver consistently outstanding results.

Requirements

In this role, you will:

- Elevate our retrieval-augmented generation (RAG) capabilities by ensuring efficient integration and utilization of diverse data sources
- Focus on enhancing embedding generation techniques, optimizing them for both accuracy and efficiency
- Stay updated on the latest advancements in AI models, understanding their evolving capabilities and complexities of integration. You'll become the team's go-to expert, assisting end-users in selecting the most suitable models for their use cases
- Collaborate closely with both backend and frontend teams to seamlessly integrate AI models into our products and services
- Prioritize customer needs and build strong relationships by actively engaging in customer support alongside other team members

You'll excel in this role if you have:

- A Bachelor's or Master's degree in Computer Science, Machine Learning, AI, or related fields
- Demonstrated expertise in machine learning and AI, particularly in generative models
- Proficiency in implementing retrieval-augmented generation pipelines
- Strong programming skills in Python and hands-on experience with popular machine learning frameworks like TensorFlow or PyTorch
- Familiarity with embedding techniques and their practical applications in AI models
- Exceptional problem-solving abilities and a track record of thriving in fast-paced environments

- Excellent communication and collaboration skills

You'll stand out if you also possess:

- A PhD in a relevant field, showcasing your advanced expertise
- Published research or projects in generative AI, embeddings, or retrieval-augmented generation
- Experience with deploying large-scale ML models in production environments
- Contributions to open-source AI/ML projects, highlighting your proactive involvement in the community
- In-depth knowledge of generative AI hardware, including GPU architectures, requirements, and optimization techniques for AI model deployment

Benefits

- Competitive salary package with performance-based bonuses
- Stock options, allowing you to share in the company's success and growth
- Comprehensive health insurance coverage for you and your dependents
- Flexible work hours and the option for remote work, promoting work-life balance
- Professional development opportunities, including access to conferences, workshops, and courses to enhance your skills
- Wellness programs and fitness stipends to support your physical and mental well-being
- Generous vacation and paid time off policies to encourage relaxation and rejuvenation
- Company-sponsored social events and team-building activities to foster a positive work culture
- Opportunities for career advancement and growth within the organization